

# FIT1047 – Week 4

## I/O and Interrupts

# Overview

Computers are almost useless without input/output.

- How does the CPU communicate with I/O devices?
- How does it handle time critical I/O?

# Early I/O

The first computers had limited I/O:

- Punched paper tape or cards
- Teleprinters

# Early I/O

The first computers had limited I/O:

- Punched paper tape or cards
- Teleprinters



# Modern I/O

# Modern I/O

- Keyboard, mouse, touch pad, touch screen, voice control, gestures, accelerometers, barometers, GPS

# Modern I/O

- Keyboard, mouse, touch pad, touch screen, voice control, gestures, accelerometers, barometers, GPS
- Screens, printers, audio, robots, ...

# Modern I/O

- Keyboard, mouse, touch pad, touch screen, voice control, gestures, accelerometers, barometers, GPS
- Screens, printers, audio, robots, ...

Also classed as I/O:



# Modern I/O

- Keyboard, mouse, touch pad, touch screen, voice control, gestures, accelerometers, barometers, GPS
- Screens, printers, audio, robots, ...

Also classed as I/O:

- External storage, network devices (WiFi, 4G, Ethernet)

# Modern I/O interfaces

I/O devices are now usually connected via interfaces:

- Standardised connectors and protocol
- Can be internal or external
- E.g. USB, SATA, HDMI, PCI Express

# I/O and the CPU

The CPU needs to

- read data from an I/O device
- write data to an I/O device

# I/O and the CPU

The CPU needs to

- read data from an I/O device
- write data to an I/O device

Why write to input devices?

- E.g. set sensitivity, calibrate, ...

# I/O and the CPU

The CPU needs to

- read data from an I/O device
- write data to an I/O device

Why write to input devices?

- E.g. set sensitivity, calibrate, ...

Why read from output devices?

- E.g. check if ready for output, check if successful, ...

# I/O and the CPU

I/O devices have their own registers.

Two ways to communicate:

Memory-mapped:

I/O registers are mapped into CPU address space.

Use Load, Store etc to communicate with I/O.

Instruction-based:

CPU has special I/O instructions.

Similar to Load, Store etc but with separate address space.

# When to do I/O

Now we know how to communicate with I/O devices.

But most I/O devices are much, much slower than the CPU. So when should the CPU communicate?

# When to do I/O

Now we know how to communicate with I/O devices.

But most I/O devices are much, much slower than the CPU. So when should the CPU communicate?

- How does it know that a new character is available from the keyboard?



# When to do I/O

Now we know how to communicate with I/O devices.

But most I/O devices are much, much slower than the CPU. So when should the CPU communicate?

- How does it know that a new character is available from the keyboard?
- How does it know a network device is ready for sending?

# Programmed I/O

Program checks registers periodically.

Also called polling I/O.

# Programmed I/O

Program checks registers periodically.

Also called polling I/O.

Pseudocode:

```
while (true) {  
    if (IORegister1.canRead()) {  
        processRegister1();  
    } else if (IORegister2.canRead()) {  
        processRegister2();  
    } else if (IORegister3.canWrite()) {  
        processRegister3();  
    }  
}
```

# Programmed I/O

## Advantage:

- We can decide how often to poll a device (prioritisation!)

## Disadvantage:

- Program is I/O-driven
- CPU is constantly in "busy loop"

Programmed I/O is mostly used in embedded special-purpose systems.

# Interrupts

Opposite of polling:

- CPU is notified when I/O is available.
- CPU interrupts what it's doing, processes I/O, then continues normal program.

# Interrupt signals

Device notifies CPU of pending interrupt by setting a bit in a special register.

CPU checks before each fetch-decode-execute cycle:

- Is interrupt bit set? Process interrupt.
- Otherwise: fetch-decode-execute.

# Interrupt signals

RTL for fetch cycle with interrupts:

```
if InterruptBit is set:  
    Clear InterruptBit  
    MAR    <- SavePC  
    M[MAR] <- PC  
    PC     <- InterruptHandler  
  
MAR    <- PC  
IR     <- M[MAR]  
PC     <- PC+1
```

The interrupt handler code is always stored at a fixed address InterruptHandler.

# Interrupt handler

Must leave the CPU in the exact same state!

- For MARIE, contents of AC must be the same as before the interrupt.

Can be achieved by

- shadow registers, a separate register file that the CPU uses while processing interrupts; or
- saving all registers to memory



# Interrupt vectors

So how can we process interrupts from different devices?

- Each device is assigned an identification number
- When raising an interrupt, stores that number (in special register or in memory)
- Interrupt handler jumps into an interrupt vector

# Interrupt vectors

```
B00      Store ACInt
B01      Load  IVec
B02      Add    DeviceID
B03      Store Dest
B04      JumpI  Dest
```


```
C00 ACInt,  HEX 0      / Temporary storage for AC
C01 Dest,   HEX 0      / Destination of jump
C02 IVec,   HEX C03
C03        HEX 0F0     / Address of first handler
C04        HEX 0FA     / Address of second handler
C05        HEX 1B0     / Address of third handler
```

# Interrupts in x86 PCs

Original design:


- 15 interrupt request (IRQ) signals
- hardware must be configured to use correct IRQ
- e.g. setting jumper on a network or sound card
- devices have to share IRQs

# Interrupts in x86 PCs





# Interrupts in x86 PCs



# Modern Interrupts

Use Advanced Programmable Interrupt Controllers (APICs).

- nowadays integrated into CPUs
- allow more IRQs (fewer conflicts)
- include high-resolution timers

# Disadvantages

- Different I/O devices need different priorities
  - can be achieved using different interrupt signals
- All memory transfers run through the CPU:
  - e.g. reads word from disk storage, writes word to memory
  - reads word from memory, writes word to graphics card
- I/O devices are fully controlled by CPU

# DMA

(Direct Memory Access)

CPU can delegate memory transfer operations to dedicated controller.

- hard disk controller can transfer data to memory
- graphics card can fetch image from memory
- while CPU can perform other tasks

CPU and DMA controller share the data bus:

- only one can do memory transfers at the same time



# Summary

## I/O

- memory-mapped vs. instruction-based
- programmed vs. interrupt-driven

## Interrupts

- require context switch
- jump into interrupt vector

## DMA

- off-load responsibility for data transfers to special controller
- keeps CPU free to do more interesting tasks

# Next lecture

- Booting up a computer
- BIOS, UEFI